

DYLAN YANG ZOU FINAL

February 21, 2024

```
[4]: import os
      os.getcwd()
      os.chdir("C:\\Data")
      os.getcwd()
```

```
[4]: 'C:\\Data'
```

```
[5]: pip install skimpy
```

```
Requirement already satisfied: skimpy in c:\users\dyland\anaconda3\lib\site-
packages (0.0.5)
Requirement already satisfied: click==7.1.2 in
c:\users\dyland\anaconda3\lib\site-packages (from skimpy) (7.1.2)
Requirement already satisfied: pandas<2.0.0,>=1.3.2 in
c:\users\dyland\anaconda3\lib\site-packages (from skimpy) (1.3.4)
Requirement already satisfied: typeguard<3.0.0,>=2.12.1 in
c:\users\dyland\anaconda3\lib\site-packages (from skimpy) (2.13.3)
Requirement already satisfied: Pygments<3.0.0,>=2.10.0 in
c:\users\dyland\anaconda3\lib\site-packages (from skimpy) (2.10.0)
Requirement already satisfied: rich<11.0.0,>=10.9.0 in
c:\users\dyland\anaconda3\lib\site-packages (from skimpy) (10.16.2)
Requirement already satisfied: pytz>=2017.3 in
c:\users\dyland\anaconda3\lib\site-packages (from pandas<2.0.0,>=1.3.2->skimpy)
(2021.3)
Requirement already satisfied: numpy>=1.17.3 in
c:\users\dyland\anaconda3\lib\site-packages (from pandas<2.0.0,>=1.3.2->skimpy)
(1.20.3)
Requirement already satisfied: python-dateutil>=2.7.3 in
c:\users\dyland\anaconda3\lib\site-packages (from pandas<2.0.0,>=1.3.2->skimpy)
(2.8.2)
Requirement already satisfied: six>=1.5 in c:\users\dyland\anaconda3\lib\site-
packages (from python-dateutil>=2.7.3->pandas<2.0.0,>=1.3.2->skimpy) (1.16.0)
Requirement already satisfied: colorama<0.5.0,>=0.4.0 in
c:\users\dyland\anaconda3\lib\site-packages (from rich<11.0.0,>=10.9.0->skimpy)
(0.4.4)
Requirement already satisfied: commonmark<0.10.0,>=0.9.0 in
c:\users\dyland\anaconda3\lib\site-packages (from rich<11.0.0,>=10.9.0->skimpy)
(0.9.1)
```

Note: you may need to restart the kernel to use updated packages.

```
[6]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import statsmodels.api as sm
import statsmodels.formula.api as smf
from patsy import dmatrix
from pyearth import Earth
from sklearn.metrics import (mean_squared_error, r2_score, roc_curve, auc,
    precision_recall_curve, make_scorer,
    recall_score, accuracy_score, precision_score,
    confusion_matrix)
from sklearn.model_selection import (cross_val_score, train_test_split, KFold,
    StratifiedKFold,
    GridSearchCV, ParameterGrid)
from sklearn.tree import DecisionTreeRegressor, DecisionTreeClassifier
from sklearn.ensemble import (BaggingRegressor, BaggingClassifier,
    RandomForestRegressor, RandomForestClassifier,
    GradientBoostingRegressor,
    GradientBoostingClassifier, AdaBoostRegressor, AdaBoostClassifier,
    VotingRegressor, VotingClassifier,
    StackingRegressor, StackingClassifier)
from sklearn.linear_model import LinearRegression, LogisticRegression, LassoCV,
    RidgeCV, ElasticNetCV
from sklearn.neighbors import KNeighborsRegressor
import itertools as it
import xgboost as xgb
import time as time
import random
from skimpy import clean_columns

#Libraries for visualizing trees
from sklearn.tree import export_graphviz
from six import StringIO
from IPython.display import Image
import pydotplus
```

1 Importing and Cleaning Data

```
[7]: data = pd.read_csv('heart_2020_cleaned.csv')
data
```

```
[7]:      HeartDisease      BMI Smoking AlcoholDrinking Stroke PhysicalHealth \
0                No  16.60      Yes                No      No                3.0
```

1	No	20.34	No	No	Yes	0.0
2	No	26.58	Yes	No	No	20.0
3	No	24.21	No	No	No	0.0
4	No	23.71	No	No	No	28.0
...
319790	Yes	27.41	Yes	No	No	7.0
319791	No	29.84	Yes	No	No	0.0
319792	No	24.24	No	No	No	0.0
319793	No	32.81	No	No	No	0.0
319794	No	46.56	No	No	No	0.0

	MentalHealth	DiffWalking	Sex	AgeCategory	Race	Diabetic	\
0	30.0	No	Female	55-59	White	Yes	
1	0.0	No	Female	80 or older	White	No	
2	30.0	No	Male	65-69	White	Yes	
3	0.0	No	Female	75-79	White	No	
4	0.0	Yes	Female	40-44	White	No	
...	
319790	0.0	Yes	Male	60-64	Hispanic	Yes	
319791	0.0	No	Male	35-39	Hispanic	No	
319792	0.0	No	Female	45-49	Hispanic	No	
319793	0.0	No	Female	25-29	Hispanic	No	
319794	0.0	No	Female	80 or older	Hispanic	No	

	PhysicalActivity	GenHealth	SleepTime	Asthma	KidneyDisease	SkinCancer
0	Yes	Very good	5.0	Yes	No	Yes
1	Yes	Very good	7.0	No	No	No
2	Yes	Fair	8.0	Yes	No	No
3	No	Good	6.0	No	No	Yes
4	Yes	Very good	8.0	No	No	No
...
319790	No	Fair	6.0	Yes	No	No
319791	Yes	Very good	5.0	Yes	No	No
319792	Yes	Good	6.0	No	No	No
319793	No	Good	12.0	No	No	No
319794	Yes	Good	8.0	No	No	No

[319795 rows x 18 columns]

```
[8]: print(data.isnull().values.any())
data.dtypes
```

False

```
[8]: HeartDisease    object
     BMI             float64
     Smoking         object
```

```

AlcoholDrinking    object
Stroke             object
PhysicalHealth     float64
MentalHealth       float64
DiffWalking        object
Sex                object
AgeCategory        object
Race               object
Diabetic           object
PhysicalActivity    object
GenHealth          object
SleepTime          float64
Asthma             object
KidneyDisease      object
SkinCancer         object
dtype: object

```

```

[9]: #convert age to numerical
import re
ages = list(data['AgeCategory'].unique())
ages = [re.findall(r'\d+', i) for i in ages]
ages = [[int(j) for j in i] for i in ages]
mean_ages = [np.array(i).mean() for i in ages]
ranges = data['AgeCategory'].unique()
d = {}
for i in range(len(mean_ages)):
    d[ranges[i]] = mean_ages[i]
print(d)
data['Age'] = data['AgeCategory'].apply(lambda x: d[x])

```

```

{'55-59': 57.0, '80 or older': 80.0, '65-69': 67.0, '75-79': 77.0, '40-44': 42.0, '70-74': 72.0, '60-64': 62.0, '50-54': 52.0, '45-49': 47.0, '18-24': 21.0, '35-39': 37.0, '30-34': 32.0, '25-29': 27.0}

```

```

[10]: data.head()

```

```

[10]:   HeartDisease    BMI Smoking AlcoholDrinking Stroke  PhysicalHealth \
0           No  16.60     Yes                No     No              3.0
1           No  20.34     No                 No     Yes              0.0
2           No  26.58     Yes                No     No             20.0
3           No  24.21     No                 No     No              0.0
4           No  23.71     No                 No     No             28.0

      MentalHealth DiffWalking    Sex AgeCategory  Race Diabetic \
0           30.0          No  Female      55-59  White     Yes
1           0.0          No  Female  80 or older  White     No
2           30.0          No   Male     65-69  White     Yes

```

3	0.0	No	Female	75-79	White	No
4	0.0	Yes	Female	40-44	White	No

	PhysicalActivity	GenHealth	SleepTime	Asthma	KidneyDisease	SkinCancer	Age
0	Yes	Very good	5.0	Yes	No	Yes	57.0
1	Yes	Very good	7.0	No	No	No	80.0
2	Yes	Fair	8.0	Yes	No	No	67.0
3	No	Good	6.0	No	No	Yes	77.0
4	Yes	Very good	8.0	No	No	No	42.0

```
[11]: pd.get_dummies(data).corr()['HeartDisease_Yes'].sort_values()
```

```
[11]: HeartDisease_No          -1.000000
DiffWalking_No              -0.201258
Stroke_No                   -0.196835
Diabetic_No                 -0.170977
KidneyDisease_No            -0.145197
GenHealth_Excellent         -0.116042
Smoking_No                  -0.107764
GenHealth_Very good         -0.101886
PhysicalActivity_Yes         -0.100030
SkinCancer_No               -0.093317
AgeCategory_18-24           -0.075385
Sex_Female                  -0.070040
AgeCategory_35-39           -0.066685
AgeCategory_25-29           -0.065759
AgeCategory_30-34           -0.065611
AgeCategory_40-44           -0.059196
AgeCategory_45-49           -0.049733
Asthma_No                   -0.041444
Race_Hispanic               -0.036163
AgeCategory_50-54           -0.032648
AlcoholDrinking_Yes         -0.032080
Race_Asian                  -0.030262
Diabetic_Yes (during pregnancy) -0.013930
AgeCategory_55-59           -0.013276
Race_Black                  -0.010156
Race_Other                  -0.003039
SleepTime                   0.008327
Race_American Indian/Alaskan Native 0.008547
AgeCategory_60-64           0.016152
Diabetic_No, borderline diabetes 0.016182
MentalHealth                0.028591
AlcoholDrinking_No          0.032080
GenHealth_Good              0.039033
Race_White                  0.040121
Asthma_Yes                  0.041444
```

AgeCategory_65-69	0.042626
BMI	0.051803
Sex_Male	0.070040
AgeCategory_70-74	0.082578
SkinCancer_Yes	0.093317
AgeCategory_75-79	0.098690
PhysicalActivity_No	0.100030
Smoking_Yes	0.107764
AgeCategory_80 or older	0.143041
KidneyDisease_Yes	0.145197
GenHealth_Fair	0.147954
PhysicalHealth	0.170721
GenHealth_Poor	0.174662
Diabetic_Yes	0.183072
Stroke_Yes	0.196835
DiffWalking_Yes	0.201258
Age	0.231583
HeartDisease_Yes	1.000000

Name: HeartDisease_Yes, dtype: float64

```
[12]: data_dum = data.drop(columns = ['SleepTime', 'AgeCategory'])
      data_dum = pd.get_dummies(data_dum)
      data_dum.columns
```

```
[12]: Index(['BMI', 'PhysicalHealth', 'MentalHealth', 'Age', 'HeartDisease_No',
            'HeartDisease_Yes', 'Smoking_No', 'Smoking_Yes', 'AlcoholDrinking_No',
            'AlcoholDrinking_Yes', 'Stroke_No', 'Stroke_Yes', 'DiffWalking_No',
            'DiffWalking_Yes', 'Sex_Female', 'Sex_Male',
            'Race_American Indian/Alaskan Native', 'Race_Asian', 'Race_Black',
            'Race_Hispanic', 'Race_Other', 'Race_White', 'Diabetic_No',
            'Diabetic_No, borderline diabetes', 'Diabetic_Yes',
            'Diabetic_Yes (during pregnancy)', 'PhysicalActivity_No',
            'PhysicalActivity_Yes', 'GenHealth_Excellent', 'GenHealth_Fair',
            'GenHealth_Good', 'GenHealth_Poor', 'GenHealth_Very good', 'Asthma_No',
            'Asthma_Yes', 'KidneyDisease_No', 'KidneyDisease_Yes', 'SkinCancer_No',
            'SkinCancer_Yes'],
            dtype='object')
```

```
[13]: no_var = []

      for col in list(pd.get_dummies(data_dum).columns):
          if col.endswith('_No'):
              no_var.append(col)

      no_var
```

```
[13]: ['HeartDisease_No',
       'Smoking_No',
       'AlcoholDrinking_No',
       'Stroke_No',
       'DiffWalking_No',
       'Diabetic_No',
       'PhysicalActivity_No',
       'Asthma_No',
       'KidneyDisease_No',
       'SkinCancer_No']
```

```
[14]: #reducing levels of categorical predictors
data_dum = data_dum.drop(columns = no_var)
data_dum = data_dum.drop(columns = ['Sex_Female', 'Race_American Indian/Alaskan',
    ↪Native',
                                'Race_Asian', 'Race_Black',
    ↪'Race_Hispanic', 'Race_Other',
                                'Diabetic_No, borderline diabetes',
    ↪'Diabetic_Yes (during pregnancy)'])
data_dum.head()
```

```
[14]:      BMI  PhysicalHealth  MentalHealth  Age  HeartDisease_Yes  Smoking_Yes  \
0   16.60             3.0           30.0  57.0                0            1
1   20.34             0.0            0.0  80.0                0            0
2   26.58            20.0           30.0  67.0                0            1
3   24.21             0.0            0.0  77.0                0            0
4   23.71            28.0           0.0  42.0                0            0

      AlcoholDrinking_Yes  Stroke_Yes  DiffWalking_Yes  Sex_Male  ...  \
0                0            0                0            0  ...
1                0            1                0            0  ...
2                0            0                0            1  ...
3                0            0                0            0  ...
4                0            0                1            0  ...

      Diabetic_Yes  PhysicalActivity_Yes  GenHealth_Excellent  GenHealth_Fair  \
0                1                1                0            0
1                0                1                0            0
2                1                1                0            1
3                0                0                0            0
4                0                1                0            0

      GenHealth_Good  GenHealth_Poor  GenHealth_Very good  Asthma_Yes  \
0                0                0                1            1
1                0                0                1            0
2                0                0                0            1
3                1                0                0            0
```

	4	0	0	1	0
		KidneyDisease_Yes	SkinCancer_Yes		
0		0	1		
1		0	0		
2		0	0		
3		0	1		
4		0	0		

[5 rows x 21 columns]

```
[15]: #didn't end up using skimpy since that made everything lowercase
data_dum.columns = data_dum.columns.str.replace(' ', '_')
data_dum.columns
```

```
[15]: Index(['BMI', 'PhysicalHealth', 'MentalHealth', 'Age', 'HeartDisease_Yes',
        'Smoking_Yes', 'AlcoholDrinking_Yes', 'Stroke_Yes', 'DiffWalking_Yes',
        'Sex_Male', 'Race_White', 'Diabetic_Yes', 'PhysicalActivity_Yes',
        'GenHealth_Excellent', 'GenHealth_Fair', 'GenHealth_Good',
        'GenHealth_Poor', 'GenHealth_Very_good', 'Asthma_Yes',
        'KidneyDisease_Yes', 'SkinCancer_Yes'],
        dtype='object')
```

```
[16]: #Balancing positive and negative responses
yes = data_dum[data_dum['HeartDisease_Yes'] == 1]
print(yes.shape)
yes.head()
```

(27373, 21)

```
[16]:      BMI  PhysicalHealth  MentalHealth  Age  HeartDisease_Yes  Smoking_Yes  \
5    28.87             6.0           0.0  77.0                1             1
10   34.30            30.0           0.0  62.0                1             1
35   32.98            10.0           0.0  77.0                1             1
42   25.06             0.0           0.0  80.0                1             0
43   30.23             6.0           2.0  77.0                1             1

      AlcoholDrinking_Yes  Stroke_Yes  DiffWalking_Yes  Sex_Male  ...  \
5                      0           0                1          0  ...
10                     0           0                1          1  ...
35                     0           1                1          1  ...
42                     0           0                1          0  ...
43                     0           0                1          0  ...

      Diabetic_Yes  PhysicalActivity_Yes  GenHealth_Excellent  GenHealth_Fair  \
5                0                   0                    0             1
10               1                   0                    0             0
```


35	1	1	0	0
42	1	0	0	0
43	1	1	0	1

	GenHealth_Good	GenHealth_Poor	GenHealth_Very_good	Asthma_Yes	\
5	0	0	0	0	
10	0	1	0	1	
35	0	1	0	0	
42	1	0	0	0	
43	0	0	0	0	

	KidneyDisease_Yes	SkinCancer_Yes
5	0	0
10	0	0
35	0	1
42	0	1
43	1	0

[5 rows x 21 columns]

```
[17]: no = data_dum[data_dum['HeartDisease_Yes'] == 0]
print(no.shape)
no = no.sample(n = 27373)
print(no.shape)
no.head()
```

(292422, 21)

(27373, 21)

```
[17]: BMI PhysicalHealth MentalHealth Age HeartDisease_Yes \
241662 36.61 0.0 30.0 47.0 0
119618 27.89 0.0 0.0 42.0 0
5352 27.44 0.0 3.0 67.0 0
29156 19.97 3.0 0.0 32.0 0
243296 45.19 0.0 15.0 32.0 0
```

	Smoking_Yes	AlcoholDrinking_Yes	Stroke_Yes	DiffWalking_Yes	\
241662	1	0	0	0	
119618	1	0	0	0	
5352	1	0	0	0	
29156	0	0	0	1	
243296	1	0	0	1	

	Sex_Male	...	Diabetic_Yes	PhysicalActivity_Yes	\
241662	0	...	0	1	
119618	1	...	0	1	
5352	0	...	0	1	

29156	0	...	0	0
243296	0	...	0	1

	GenHealth_Excellent	GenHealth_Fair	GenHealth_Good	GenHealth_Poor	\
241662	0	0	1	0	
119618	0	0	1	0	
5352	0	0	1	0	
29156	0	1	0	0	
243296	0	0	1	0	

	GenHealth_Very_good	Asthma_Yes	KidneyDisease_Yes	SkinCancer_Yes
241662	0	0	0	0
119618	0	0	0	0
5352	0	0	0	0
29156	0	1	0	0
243296	0	1	0	0

[5 rows x 21 columns]

```
[18]: data_balanced = pd.concat([no, yes], ignore_index = True)
      print(data_balanced.shape)
      data_balanced.head()
```

(54746, 21)

```
[18]: BMI    PhysicalHealth    MentalHealth    Age    HeartDisease_Yes    Smoking_Yes    \
0    36.61            0.0            30.0    47.0            0            1
1    27.89            0.0            0.0    42.0            0            1
2    27.44            0.0            3.0    67.0            0            1
3    19.97            3.0            0.0    32.0            0            0
4    45.19            0.0            15.0    32.0            0            1
```

	AlcoholDrinking_Yes	Stroke_Yes	DiffWalking_Yes	Sex_Male	...	\
0	0	0	0	0	...	
1	0	0	0	1	...	
2	0	0	0	0	...	
3	0	0	1	0	...	
4	0	0	1	0	...	

	Diabetic_Yes	PhysicalActivity_Yes	GenHealth_Excellent	GenHealth_Fair	\
0	0	1	0	0	
1	0	1	0	0	
2	0	1	0	0	
3	0	0	0	1	
4	0	1	0	0	

	GenHealth_Good	GenHealth_Poor	GenHealth_Very_good	Asthma_Yes	\
--	----------------	----------------	---------------------	------------	---

0	1	0	0	0
1	1	0	0	0
2	1	0	0	0
3	0	0	0	1
4	1	0	0	1

	KidneyDisease_Yes	SkinCancer_Yes
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

[5 rows x 21 columns]

```
[19]: X = data_balanced.drop(columns = 'HeartDisease_Yes')
X.head()
```

```
[19]:      BMI  PhysicalHealth  MentalHealth  Age  Smoking_Yes  \
0  36.61             0.0           30.0  47.0             1
1  27.89             0.0             0.0  42.0             1
2  27.44             0.0             3.0  67.0             1
3  19.97             3.0             0.0  32.0             0
4  45.19             0.0           15.0  32.0             1
```

	AlcoholDrinking_Yes	Stroke_Yes	DiffWalking_Yes	Sex_Male	Race_White	\
0	0	0	0	0	1	
1	0	0	0	1	0	
2	0	0	0	0	0	
3	0	0	1	0	1	
4	0	0	1	0	0	

	Diabetic_Yes	PhysicalActivity_Yes	GenHealth_Excellent	GenHealth_Fair	\
0	0	1	0	0	
1	0	1	0	0	
2	0	1	0	0	
3	0	0	0	1	
4	0	1	0	0	

	GenHealth_Good	GenHealth_Poor	GenHealth_Very_good	Asthma_Yes	\
0	1	0	0	0	
1	1	0	0	0	
2	1	0	0	0	
3	0	0	0	1	
4	1	0	0	1	

	KidneyDisease_Yes	SkinCancer_Yes
--	-------------------	----------------

```

0          0          0
1          0          0
2          0          0
3          0          0
4          0          0

```

```
[20]: y = data_balanced.HeartDisease_Yes
      y.head()
```

```
[20]: 0    0
      1    0
      2    0
      3    0
      4    0
      Name: HeartDisease_Yes, dtype: uint8
```

```
[21]: X_train, X_test, y_train, y_test = train_test_split(X, y, train_size = 4000,
      ↪test_size = 1000, random_state = 1, stratify = y)
      X_train2, X_test2, y_train2, y_test2 = train_test_split(X, y, train_size =
      ↪4000, test_size = 1000, random_state = 2, stratify = y)
```

```
[22]: #Checking to make sure responses are still balanced
      print(y_train.value_counts())
      print(y_test.value_counts())
      print(y_train2.value_counts())
      print(y_test2.value_counts())
```

```

0    2000
1    2000
Name: HeartDisease_Yes, dtype: int64
0    500
1    500
Name: HeartDisease_Yes, dtype: int64
0    2000
1    2000
Name: HeartDisease_Yes, dtype: int64
0    500
1    500
Name: HeartDisease_Yes, dtype: int64

```

```
[23]: #Checking if X_train and X_train2 overlap
      np.unique([X_train.index.values == X_train2.index.values], return_counts = True)
```

```
[23]: (array([False]), array([4000], dtype=int64))
```

```
[24]: X_train.head()
```

```
[24]:
```

	BMI	PhysicalHealth	MentalHealth	Age	Smoking_Yes	\
16548	30.34	0.0	0.0	67.0	0	
1305	25.83	0.0	0.0	72.0	1	
8061	23.71	0.0	0.0	77.0	0	
20173	32.93	20.0	3.0	62.0	1	
44566	24.11	0.0	0.0	77.0	1	

	AlcoholDrinking_Yes	Stroke_Yes	DiffWalking_Yes	Sex_Male	Race_White	\
16548	0	0	0	0	1	
1305	0	0	0	1	1	
8061	0	0	0	1	1	
20173	0	0	1	0	1	
44566	0	0	0	1	1	

	Diabetic_Yes	PhysicalActivity_Yes	GenHealth_Excellent	\
16548	0	1	1	
1305	0	1	1	
8061	0	1	0	
20173	1	0	0	
44566	1	1	0	

	GenHealth_Fair	GenHealth_Good	GenHealth_Poor	GenHealth_Very_good	\
16548	0	0	0	0	
1305	0	0	0	0	
8061	0	1	0	0	
20173	1	0	0	0	
44566	1	0	0	0	

	Asthma_Yes	KidneyDisease_Yes	SkinCancer_Yes
16548	0	0	0
1305	0	0	0
8061	0	0	1
20173	0	0	0
44566	0	0	0

2 Functions

```
[25]: #Function to compute confusion matrix and prediction accuracy on test/train
      ↪ data using tree models
def confusion_matrix_data(data,actual_values,model,cutoff=0.5):
    #Predict the values using the Logit model
    pred_values = model.predict_proba(data)[:,-1]
    # Specify the bins
    bins=np.array([0,cutoff,1])
    #Confusion matrix
    cm = np.histogram2d(actual_values, pred_values, bins=bins)[0]
```

```

cm_df = pd.DataFrame(cm)
cm_df.columns = ['Predicted 0', 'Predicted 1']
cm_df = cm_df.rename(index={0: 'Actual 0', 1: 'Actual 1'})
# Calculate the accuracy
accuracy = 100*(cm[0,0]+cm[1,1])/cm.sum()
fnr = 100*(cm[1,0])/(cm[1,0]+cm[1,1])
precision = 100*(cm[1,1])/(cm[0,1]+cm[1,1])
fpr = 100*(cm[0,1])/(cm[0,0]+cm[0,1])
tpr = 100*(cm[1,1])/(cm[1,0]+cm[1,1])
print("Accuracy = ", accuracy)
print("Precision = ", precision)
print("FNR = ", fnr)
# print("FPR = ", fpr)
print("TPR or Recall = ", tpr)
print("Confusion matrix = \n", cm_df)
return (" ")

```

```

[26]: def plot_precision_recall_vs_threshold(precisions, recalls, thresholds):
    plt.figure(figsize = (8, 8))
    plt.title('Precision and Recall Scores as a function of the decision_
    ↪threshold')
    plt.plot(thresholds, precisions[:-1], 'b--', label = 'Precision')
    plt.plot(thresholds, recalls[:-1], 'g-', label = 'Recall')
    plt.ylabel('Score')
    plt.xlabel('Decision Threshold')
    plt.legend(loc = 'best')

```

```

[27]: def ols_formula(df, dependent_var):
    df_columns = list(df.columns.values)
    df_columns.remove(dependent_var)
    return dependent_var + ' ~ ' + ' + '.join(df_columns)

```

```

[28]: def confusion_matrix_train(model, cutoff = 0.5):
    # Confusion matrix
    cm_df = pd.DataFrame(model.pred_table(threshold = cutoff))
    #Formatting the confusion matrix
    cm_df.columns = ['Predicted 0', 'Predicted 1']
    cm_df = cm_df.rename(index={0: 'Actual 0', 1: 'Actual 1'})
    cm = np.array(cm_df)
    # Calculate the accuracy
    accuracy = 100*(cm[0,0]+cm[1,1])/cm.sum()
    fnr = 100*(cm[1,0])/(cm[1,0]+cm[1,1])
    precision = 100*(cm[1,1])/(cm[0,1]+cm[1,1])
    fpr = 100*(cm[0,1])/(cm[0,0]+cm[0,1])
    tpr = 100*(cm[1,1])/(cm[1,0]+cm[1,1])
    print("Accuracy = ", accuracy)
    print("Precision = ", precision)

```

```

print("FNR = ", fnr)
#print("FPR = ", fpr)
print("TPR or Recall = ", tpr)
print("Confusion matrix = \n", cm_df)
return ( " ")
return cm_df, accuracy

```

```

[29]: #Function to compute confusion matrix and prediction accuracy on test data
def confusion_matrix_test(data,actual_values,model,cutoff=0.5): #Predict the
    ↪ values using the Logit model
    pred_values = model.predict(data) # Specify the bins
    bins=np.array([0,cutoff,1]) #Confusion matrix
    cm = np.histogram2d(actual_values, pred_values, bins=bins)[0]
    cm_df = pd.DataFrame(model.pred_table(threshold = cutoff))
    cm_df.columns = ['Predicted 0','Predicted 1']
    cm_df = cm_df.rename(index={0: 'Actual 0',1:'Actual 1'})
    # Calculate the accuracy
    accuracy = 100*(cm[0,0]+cm[1,1])/cm.sum() # Return the confusion matrix and
    ↪ the accuracy
    fnr = 100*(cm[1,0])/(cm[1,0]+cm[1,1])
    precision = 100*(cm[1,1])/(cm[0,1]+cm[1,1])
    fpr = 100*(cm[0,1])/(cm[0,0]+cm[0,1])
    tpr = 100*(cm[1,1])/(cm[1,0]+cm[1,1])
    print("Accuracy = ", accuracy)
    print("Precision = ", precision)
    print("FNR = ", fnr)
    #print("FPR = ", fpr)
    print("TPR or Recall = ", tpr)
    print("Confusion matrix = \n", cm_df)
    return ( " ")

```

3 Classification Tree

```

[34]: #Defining the object to build a regression tree
model = DecisionTreeClassifier(random_state=1, max_depth=3)

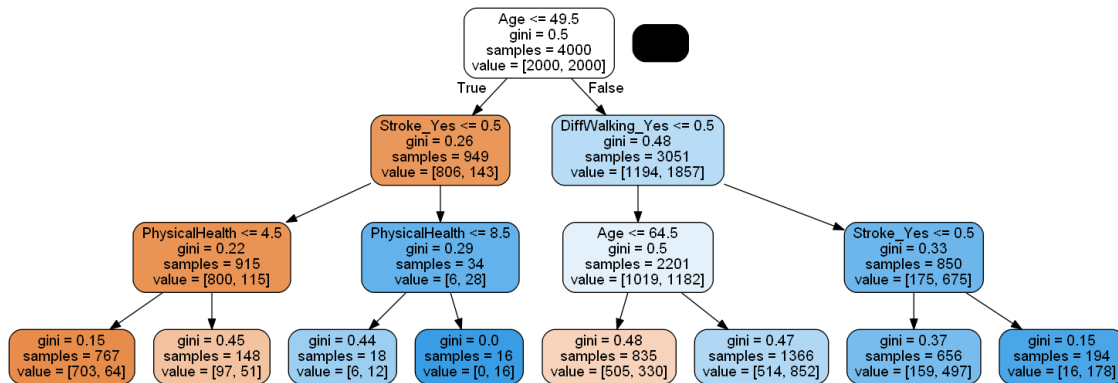
#Fitting the regression tree to the data
model.fit(X_train, y_train)

#Visualizing the regression tree
dot_data = StringIO()
export_graphviz(model, out_file=dot_data,
filled=True, rounded=True,
feature_names = X_train.columns,precision=2)
graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
#graph.write_png('car_price_tree.png')

```

```
Image(graph.create_png())
```

[34]:



```
[35]: pred=model.predict_proba(X_test)[: ,0]
confusion_matrix_data(X_train,y_train,model,cutoff=0.4)
```

```
Accuracy = 71.5
Precision = 69.11111111111111
FNR = 22.25
TPR or Recall = 77.75
Confusion matrix =
      Predicted 0 Predicted 1
Actual 0      1305.0      695.0
Actual 1       445.0     1555.0
```

[35]: ' '

```
[39]: #Defining parameters and the range of values over which to optimize
param_grid = {
    'max_depth': range(1,10),
    'max_leaf_nodes': range(10,30),
}

skf = StratifiedKFold(n_splits=5)
#The folds are made by preserving the percentage of samples for each class.
#Minimizing FNR is equivalent to maximizing recall
grid_search = GridSearchCV(DecisionTreeClassifier(random_state=1), param_grid,
                           scoring=['precision','recall','accuracy'],
                           refit="recall", cv=skf, n_jobs=-1, verbose = True)
grid_search.fit(X_train, y_train)
# make the predictions
y_pred = grid_search.predict(X_test)
print('Best params for recall')
print(grid_search.best_params_)
```


Fitting 5 folds for each of 180 candidates, totalling 900 fits
Best params for recall
{'max_depth': 2, 'max_leaf_nodes': 10}

```
[46]: #All results of the grid search can be seen with cv_results_  
cv_scores = pd.DataFrame(grid_search.cv_results_)  
  
model = DecisionTreeClassifier(random_state=1, max_depth = 5, max_leaf_nodes=25)  
model.fit(X_train,y_train)  
print(confusion_matrix_data(X_train,y_train,model))  
print(confusion_matrix_data(X_test,y_test,model))  
print(confusion_matrix_data(X_train,y_train,model,cutoff=0.2))  
print(confusion_matrix_data(X_test,y_test,model,cutoff=0.2))
```

Accuracy = 74.2
Precision = 72.83018867924528
FNR = 22.8
TPR or Recall = 77.2
Confusion matrix =

	Predicted 0	Predicted 1
Actual 0	1424.0	576.0
Actual 1	456.0	1544.0

Accuracy = 72.5
Precision = 70.49180327868852
FNR = 22.6
TPR or Recall = 77.4
Confusion matrix =

	Predicted 0	Predicted 1
Actual 0	338.0	162.0
Actual 1	113.0	387.0

Accuracy = 65.925
Precision = 59.77293648358392
FNR = 2.6
TPR or Recall = 97.4
Confusion matrix =

	Predicted 0	Predicted 1
Actual 0	689.0	1311.0
Actual 1	52.0	1948.0

Accuracy = 66.5
Precision = 60.19777503090235
FNR = 2.6
TPR or Recall = 97.4
Confusion matrix =

	Predicted 0	Predicted 1
Actual 0	178.0	322.0

Actual 1 13.0 487.0

```
[56]: clf = DecisionTreeClassifier(random_state=1)
param_grid = {
    'max_depth': range(1,10),
    'max_leaf_nodes': range(10,30),
    'max_features': range(3,8)
}

refit_score = "roc_auc"

skf = StratifiedKFold(n_splits=5)
grid_search = GridSearchCV(clf,
    ↳param_grid,scoring=['precision','recall','accuracy','roc_auc'],
                                refit=refit_score,cv=skf, n_jobs=-1, verbose = True)
grid_search.fit(X_train, y_train)

# make the predictions
y_pred = grid_search.predict(X_test)

print('Best params for', refit_score)
print(grid_search.best_params_)
```

Fitting 5 folds for each of 900 candidates, totalling 4500 fits

Best params for roc_auc

{'max_depth': 7, 'max_features': 5, 'max_leaf_nodes': 28}

```
[57]: model = DecisionTreeClassifier(random_state=1, max_depth = 7, max_leaf_nodes =
    ↳28, max_features = 5)
model.fit(X_train,y_train)
print(confusion_matrix_data(X_train,y_train,model))
print(confusion_matrix_data(X_test,y_test,model))
print(confusion_matrix_data(X_train,y_train,model,cutoff=0.2))
print(confusion_matrix_data(X_test,y_test,model,cutoff=0.2))
```

Accuracy = 75.425

Precision = 72.8744939271255

FNR = 19.0

TPR or Recall = 81.0

Confusion matrix =

	Predicted 0	Predicted 1
Actual 0	1397.0	603.0
Actual 1	380.0	1620.0

Accuracy = 74.1

Precision = 71.63375224416517

FNR = 20.2

```

TPR or Recall = 79.8
Confusion matrix =
      Predicted 0 Predicted 1
Actual 0      342.0      158.0
Actual 1      101.0      399.0

```

```

Accuracy = 66.725
Precision = 60.44985941893158
FNR = 3.25
TPR or Recall = 96.75
Confusion matrix =
      Predicted 0 Predicted 1
Actual 0      734.0      1266.0
Actual 1       65.0      1935.0

```

```

Accuracy = 67.3
Precision = 61.10397946084724
FNR = 4.8
TPR or Recall = 95.2
Confusion matrix =
      Predicted 0 Predicted 1
Actual 0      197.0      303.0
Actual 1       24.0      476.0

```

```

[36]: model = DecisionTreeClassifier(random_state = 1) #model without any restrictions
path= model.cost_complexity_pruning_path(X_train,y_train) # Compute the pruning
↳ path during Minimal Cost-Complexity Pruning.
alphas=path['ccp_alphas']
len(alphas)

```

[36]: 408

```

[58]: #Grid search to optimize parameter values
param_grid = {
    'ccp_alpha':alphas,
}

skf = StratifiedKFold(n_splits=5)
grid_search = GridSearchCV(DecisionTreeClassifier(random_state = 1), param_grid,
↳ scoring=['precision','recall','accuracy'],
refit="recall", cv=skf,
↳ n_jobs=-1, verbose = True)
grid_search.fit(X_train, y_train)

# make the predictions

```

```

y_pred = grid_search.predict(X_test)

print('Best params for recall')
print(grid_search.best_params_)

```

Fitting 5 folds for each of 408 candidates, totalling 2040 fits
 Best params for recall
 {'ccp_alpha': 0.020264403994903202}

```

[59]: tree = DecisionTreeClassifier(ccp_alpha=0.020264403994903202,random_state=1)
tree.fit(X_train, y_train)
print(confusion_matrix_data(X_train,y_train,tree))
print(confusion_matrix_data(X_test,y_pred,tree))
print(confusion_matrix_data(X_train,y_train,model,cutoff=0.2))
print(confusion_matrix_data(X_test,y_pred,model,cutoff=0.2))

```

Accuracy = 66.575
 Precision = 60.86529006882989
 FNR = 7.15
 TPR or Recall = 92.85
 Confusion matrix =

	Predicted 0	Predicted 1
Actual 0	806.0	1194.0
Actual 1	143.0	1857.0

Accuracy = 100.0
 Precision = 100.0
 FNR = 0.0
 TPR or Recall = 100.0
 Confusion matrix =

	Predicted 0	Predicted 1
Actual 0	229.0	0.0
Actual 1	0.0	771.0

Accuracy = 66.725
 Precision = 60.44985941893158
 FNR = 3.25
 TPR or Recall = 96.75
 Confusion matrix =


	Predicted 0	Predicted 1
Actual 0	734.0	1266.0
Actual 1	65.0	1935.0

Accuracy = 93.6
 Precision = 95.37869062901156
 FNR = 3.6316472114137484
 TPR or Recall = 96.36835278858625
 Confusion matrix =

	Predicted 0	Predicted 1
Actual 0	193.0	36.0
Actual 1	28.0	743.0

```
[60]: ypred = tree.predict_proba(X_train)[: , 1]
      p, r, thresholds = precision_recall_curve(y_train, ypred)


      plot_precision_recall_vs_threshold(p, r, thresholds)
```



4 Exploratory Data Analysis

```
[243]: sns.histplot(data = data, x = 'HeartDisease', shrink = 0.8)
```

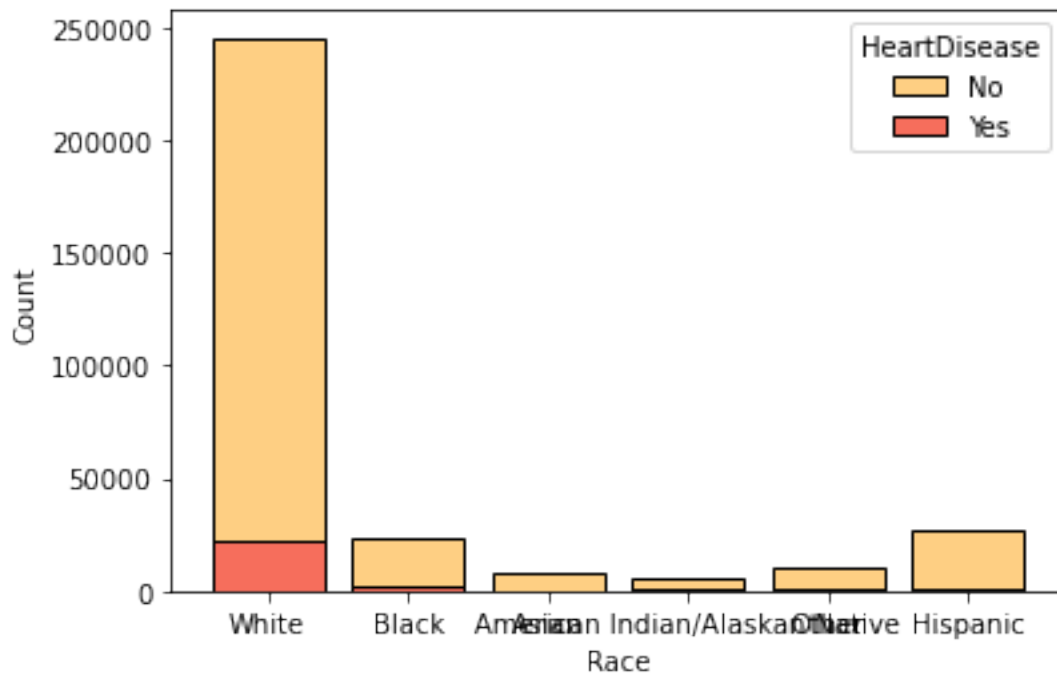
```
[243]: <AxesSubplot:xlabel='HeartDisease', ylabel='Count'>
```



- There are many more negative samples, so precision-recall is appropriate

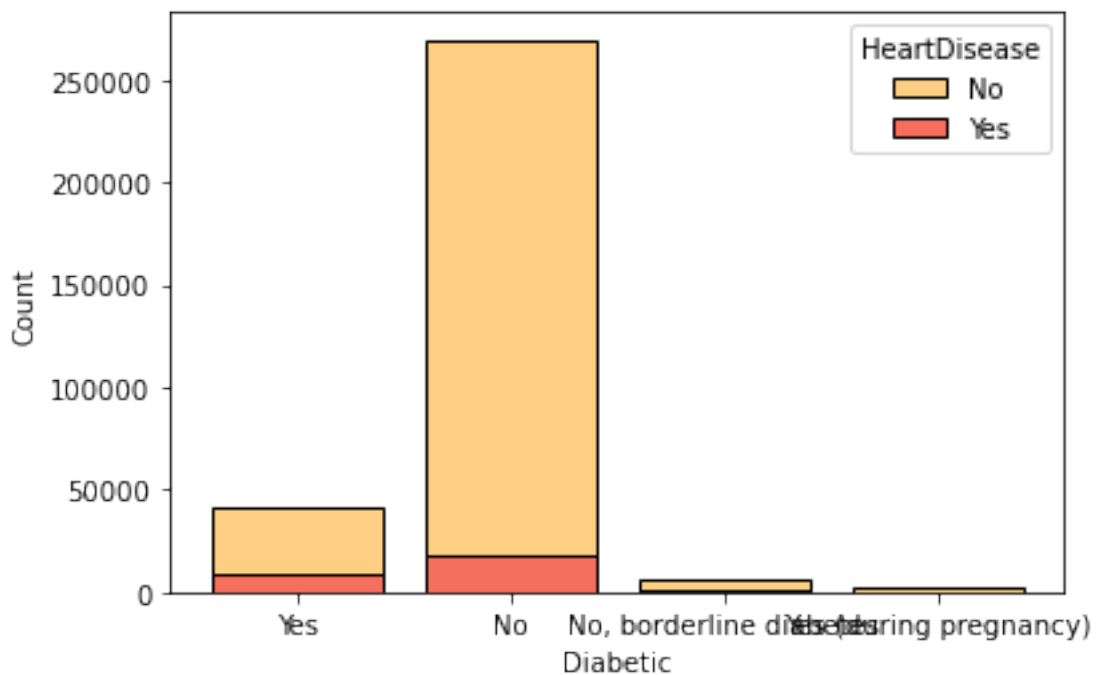
```
[89]: sns.histplot(x = "Race", hue = "HeartDisease", data = data,  
                  stat = "count", shrink = 0.8, multiple = 'stack', palette="YlOrRd")
```

```
[89]: <AxesSubplot:xlabel='Race', ylabel='Count'>
```




```
[90]: sns.histplot(x = "Diabetic", hue = "HeartDisease", data = data,
                  stat = "count", shrink = 0.8, multiple = 'stack', palette="YlOrRd")
```

```
[90]: <AxesSubplot:xlabel='Diabetic', ylabel='Count'>
```




```
[93]: sns.histplot(x = "GenHealth", hue = "HeartDisease", data = data,  
                  stat = "count", shrink = 0.8, multiple = 'stack', palette="YlOrRd")
```

```
[93]: <AxesSubplot:xlabel='GenHealth', ylabel='Count'>
```



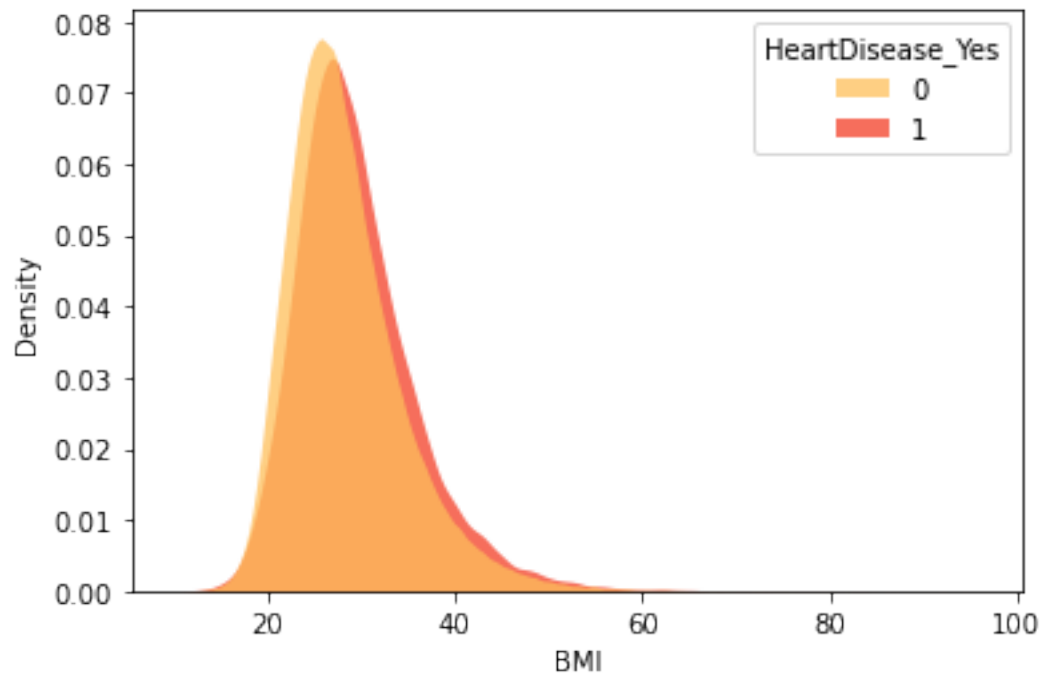
```
[153]: sns.boxplot(x="HeartDisease_Yes", y="BMI", data=data_dum, palette='YlOrRd')
```

```
[153]: <AxesSubplot:xlabel='HeartDisease_Yes', ylabel='BMI'>
```


```
[154]: sns.kdeplot(  
    data=data_dum, x="BMI", hue="HeartDisease_Yes",  
    fill=True, common_norm=False, palette="YlOrRd",  
    linewidth=0, alpha=.75  
)
```

```
[154]: <AxesSubplot:xlabel='BMI', ylabel='Density'>
```




```
[98]: sns.boxplot(x="HeartDisease_Yes", y="PhysicalHealth", data=data_dum,
→palette='YlOrRd')
```

```
[98]: <AxesSubplot:xlabel='HeartDisease_Yes', ylabel='PhysicalHealth'>
```



```
[99]: sns.boxplot(x="HeartDisease_Yes", y="MentalHealth", data=data_dum,
↳ palette='YlOrRd')
```

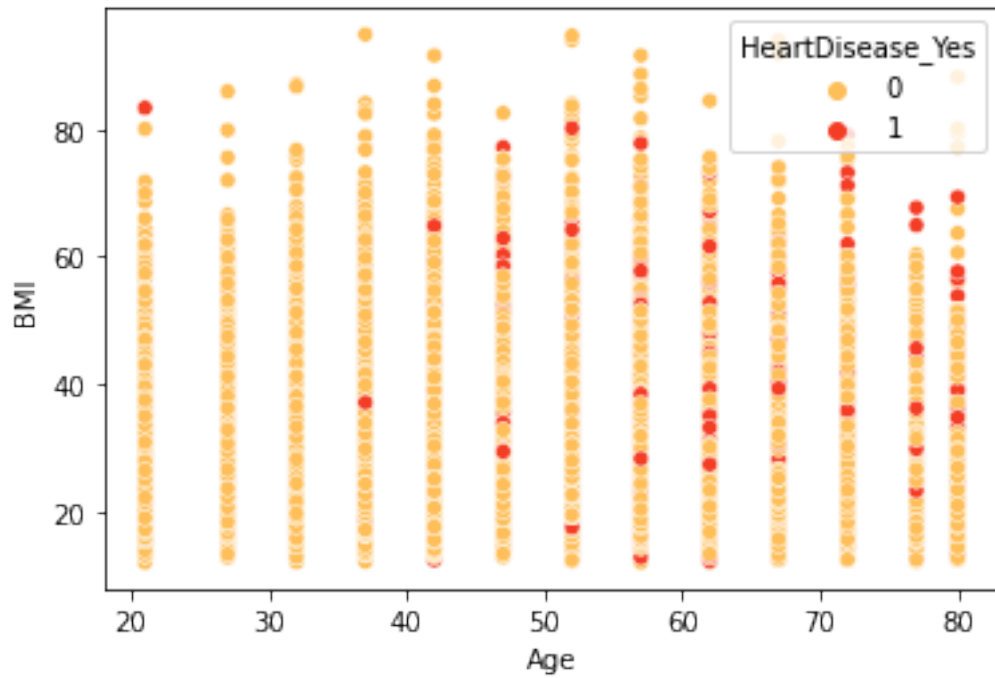
```
[99]: <AxesSubplot:xlabel='HeartDisease_Yes', ylabel='MentalHealth'>
```



```
[156]: sns.scatterplot(data=data_dum, x="Age", y="BMI", hue="HeartDisease_Yes",
↳ palette='YlOrRd')
```


```
[156]: <AxesSubplot:xlabel='Age', ylabel='BMI'>
```

```
C:\ProgramData\Anaconda3\lib\site-packages\IPython\core\pylabtools.py:132:
UserWarning: Creating legend with loc="best" can be slow with large amounts of
data.
  fig.canvas.print_figure(bytes_io, **kw)
```




```
[240]: sns.histplot(data=data, x = "Age", stat = 'count', palette = "YlOrRd")
```

```
[240]: <AxesSubplot:xlabel='Age', ylabel='Count'>
```




```
[158]: sns.barplot(data=data_dum, x="Age", y = 'HeartDisease_Yes', palette = "YlOrRd")
```

```
[158]: <AxesSubplot:xlabel='Age', ylabel='HeartDisease_Yes'>
```



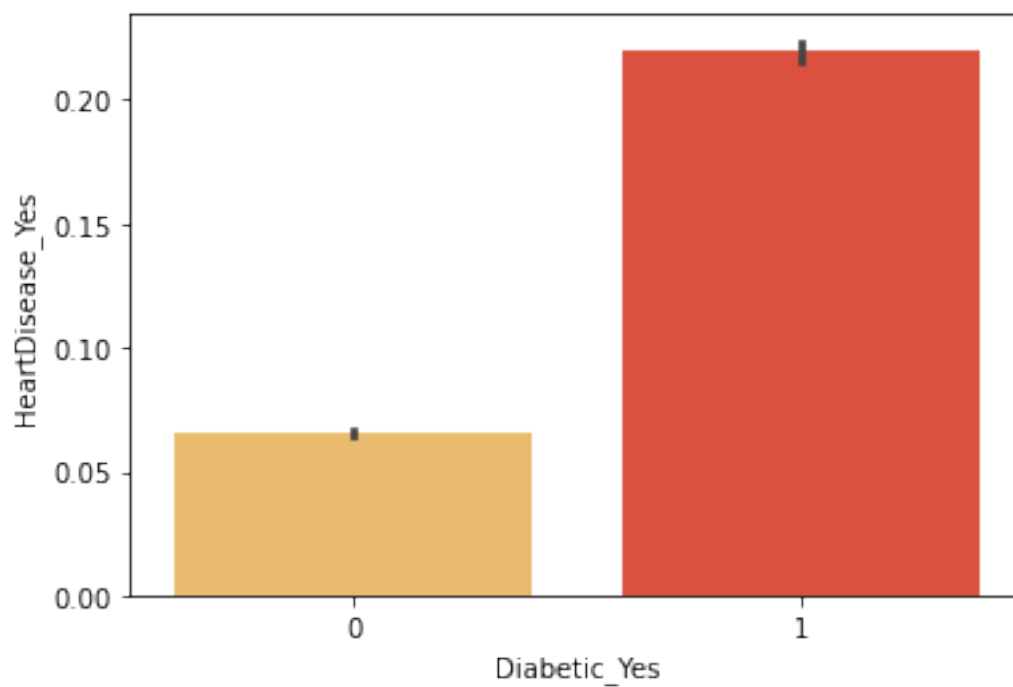
```
[159]: sns.barplot(data=data_dum, x="Stroke_Yes", y = 'HeartDisease_Yes', palette = "YlOrRd")
```

```
[159]: <AxesSubplot:xlabel='Stroke_Yes', ylabel='HeartDisease_Yes'>
```




```
[160]: sns.barplot(data=data_dum, x="Diabetic_Yes", y = 'HeartDisease_Yes', palette = "YlOrRd")
```

```
[160]: <AxesSubplot:xlabel='Diabetic_Yes', ylabel='HeartDisease_Yes'>
```




```
[100]: sns.barplot(data=data_dum, x="Smoking_Yes", y = 'HeartDisease_Yes', palette = "YlOrRd")
```

```
[100]: <AxesSubplot:xlabel='Smoking_Yes', ylabel='HeartDisease_Yes'>
```



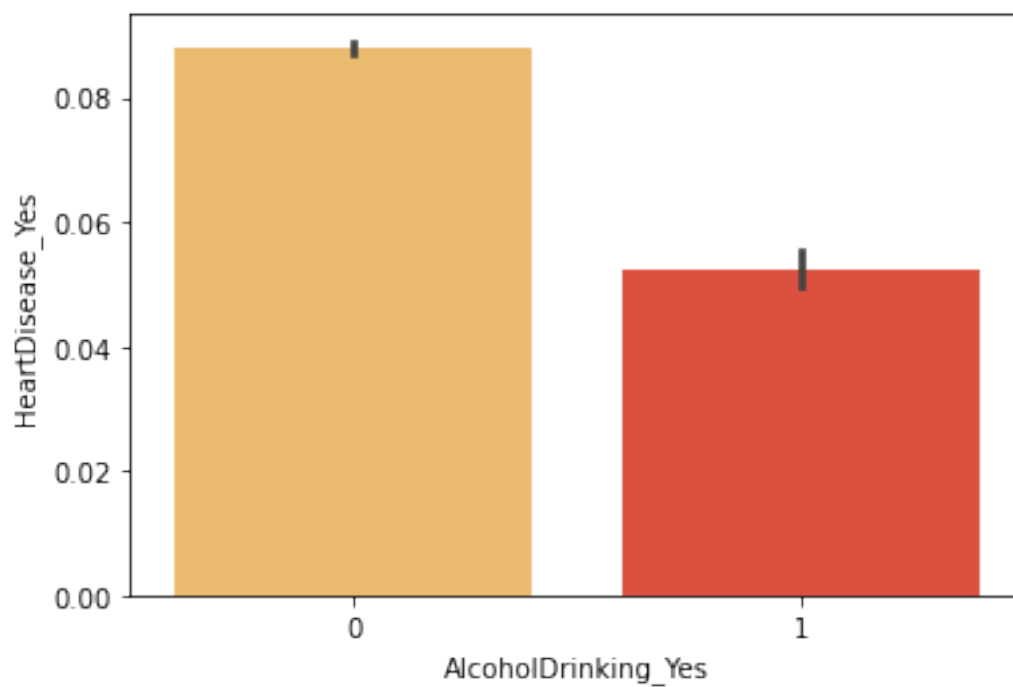
```
[149]: sns.barplot(data=data_dum, x = 'Sex_Male', y = 'HeartDisease_Yes', palette = "YlOrRd")
```

```
[149]: <AxesSubplot:xlabel='Sex_Male', ylabel='HeartDisease_Yes'>
```



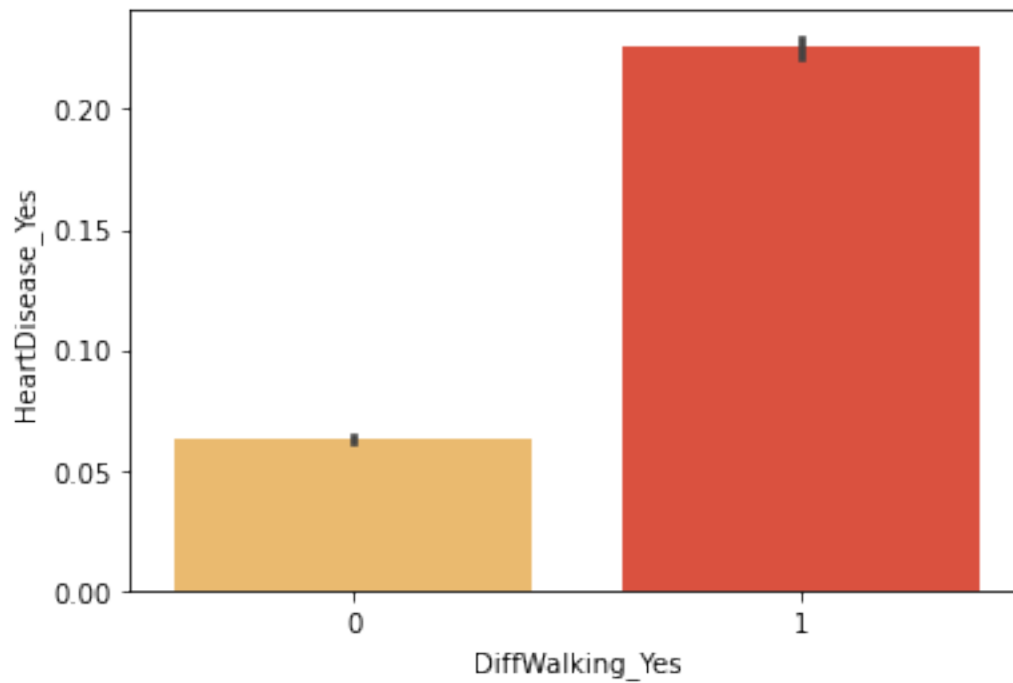
```
[150]: sns.barplot(data=data_dum, x = 'AlcoholDrinking_Yes', y = 'HeartDisease_Yes',  
→palette = "YlOrRd")
```

```
[150]: <AxesSubplot:xlabel='AlcoholDrinking_Yes', ylabel='HeartDisease_Yes'>
```





```
[151]: sns.barplot(data=data_dum, x = 'DiffWalking_Yes', y = 'HeartDisease_Yes',  
↳ palette = "YlOrRd")
```

```
[151]: <AxesSubplot:xlabel='DiffWalking_Yes', ylabel='HeartDisease_Yes'>
```




```
[152]: sns.barplot(data=data_dum, x = 'Diabetic_Yes', y = 'HeartDisease_Yes', palette=  
↳ "YlOrRd")
```

```
[152]: <AxesSubplot:xlabel='Diabetic_Yes', ylabel='HeartDisease_Yes'>
```



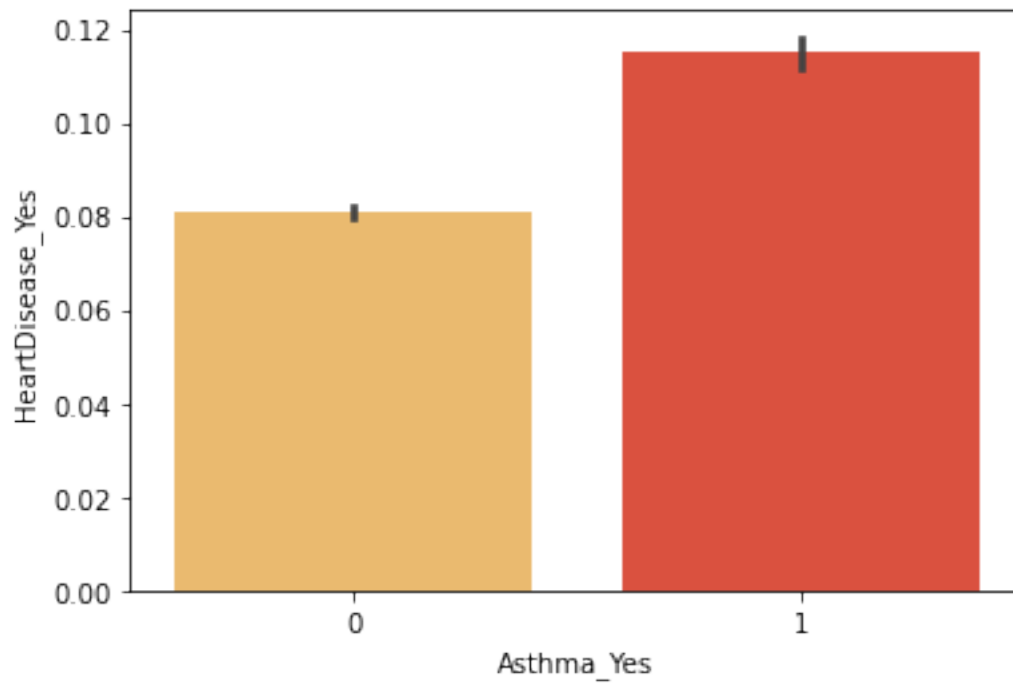
```
[153]: sns.barplot(data=data_dum, x = 'PhysicalActivity_Yes', y = 'HeartDisease_Yes',  
→palette = "YlOrRd")
```

```
[153]: <AxesSubplot:xlabel='PhysicalActivity_Yes', ylabel='HeartDisease_Yes'>
```



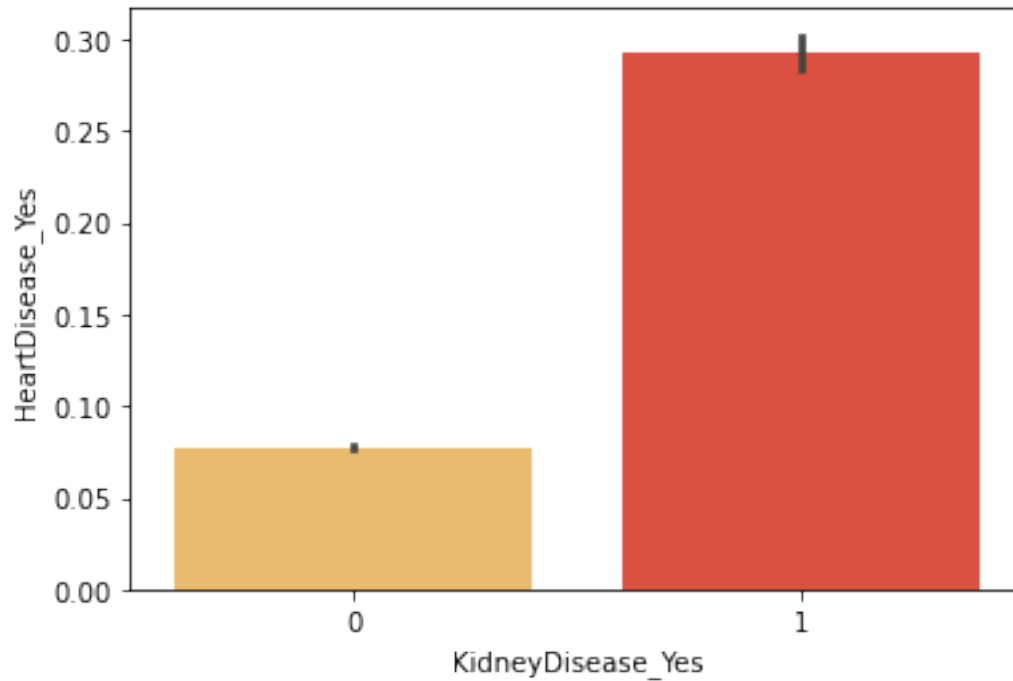
```
[154]: sns.barplot(data=data_dum, x = 'Asthma_Yes', y = 'HeartDisease_Yes', palette = "YlOrRd")
```

```
[154]: <AxesSubplot:xlabel='Asthma_Yes', ylabel='HeartDisease_Yes'>
```



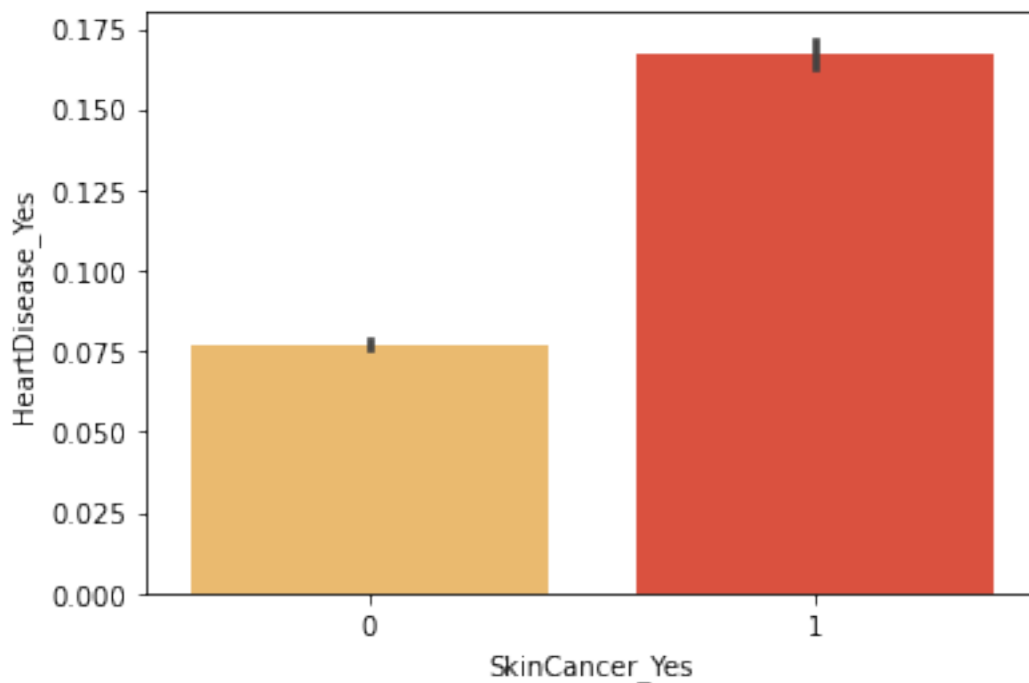
```
[156]: sns.barplot(data=data_dum, x = 'KidneyDisease_Yes', y = 'HeartDisease_Yes', palette = "YlOrRd")
```

```
[156]: <AxesSubplot:xlabel='KidneyDisease_Yes', ylabel='HeartDisease_Yes'>
```




```
[157]: sns.barplot(data=data_dum, x = 'SkinCancer_Yes', y = 'HeartDisease_Yes',
→palette = "YlOrRd")
```

```
[157]: <AxesSubplot:xlabel='SkinCancer_Yes', ylabel='HeartDisease_Yes'>
```



```
[182]: sns.barplot(data=data_dum, x = 'Race_White', y = 'HeartDisease_Yes', palette = "YlOrRd")
```

```
[182]: <AxesSubplot:xlabel='Race_White', ylabel='HeartDisease_Yes'>
```




```
[161]: fig, axes = plt.subplots(1, 1, figsize = (8,8))

p2 = sns.boxplot(data = data_dum, x = 'HeartDisease_Yes', y = 'SleepTime', palette = 'YlOrRd')

p2.set_title("")
plt.xticks([0,1],['No', 'Yes'])
```

```
[161]: ([<matplotlib.axis.XTick at 0x1ca0855fee0>,
<matplotlib.axis.XTick at 0x1ca0855feb0>],
[Text(0, 0, 'No'), Text(1, 0, 'Yes')])
```



5 Simple Logistic Regression

```
[162]: train = pd.concat([X_train, y_train], axis=1, join="inner")
train.head()
```

```
[162]:
```

	BMI	PhysicalHealth	MentalHealth	SleepTime	Smoking_Yes	\
291133	24.53	2	0	6	1	
309330	33.00	0	0	7	1	
307280	27.60	0	5	9	0	
7469	28.70	0	14	8	1	
33630	28.98	0	0	7	1	

	AlcoholDrinking_Yes	Stroke_Yes	DiffWalking_Yes	Sex_Female	\
--	---------------------	------------	-----------------	------------	---

291133	0	0	0	1
309330	0	0	0	0
307280	0	0	0	0
7469	0	0	0	0
33630	0	0	0	0

	Sex_Male	...	PhysicalActivity_Yes	GenHealth_Excellent	\
291133	0	...	1	0	
309330	1	...	1	1	
307280	1	...	1	1	
7469	1	...	1	1	
33630	1	...	1	0	

	GenHealth_Fair	GenHealth_Good	GenHealth_Poor	GenHealth_Very_good	\
291133	1	0	0	0	
309330	0	0	0	0	
307280	0	0	0	0	
7469	0	0	0	0	
33630	0	1	0	0	

	Asthma_Yes	KidneyDisease_Yes	SkinCancer_Yes	HeartDisease_Yes
291133	0	0	0	0
309330	0	0	0	0
307280	0	0	0	0
7469	0	0	0	0
33630	0	0	0	1

[5 rows x 42 columns]

```
[163]: test = pd.concat([X_test, y_test], axis=1, join="inner")
test.head()
```

	BMI	PhysicalHealth	MentalHealth	SleepTime	Smoking_Yes	\
301988	24.30	0	15	7	0	
223127	23.78	0	0	7	1	
216797	20.60	0	0	7	1	
234217	28.29	0	20	4	1	
30822	33.00	0	0	8	0	

	AlcoholDrinking_Yes	Stroke_Yes	DiffWalking_Yes	Sex_Female	\
301988	0	0	0	1	
223127	0	1	1	1	
216797	0	0	0	1	
234217	0	0	0	1	
30822	0	0	0	0	

	Sex_Male	...	PhysicalActivity_Yes	GenHealth_Excellent	\
--	----------	-----	----------------------	---------------------	---

301988	0	...	1	1
223127	0	...	0	0
216797	0	...	1	0
234217	0	...	0	0
30822	1	...	1	1

	GenHealth_Fair	GenHealth_Good	GenHealth_Poor	GenHealth_Very_good	\
301988	0	0	0	0	
223127	0	1	0	0	
216797	0	1	0	0	
234217	0	1	0	0	
30822	0	0	0	0	

	Asthma_Yes	KidneyDisease_Yes	SkinCancer_Yes	HeartDisease_Yes
301988	0	0	0	0
223127	0	0	0	1
216797	0	0	0	0
234217	0	0	0	0
30822	0	0	0	1

[5 rows x 42 columns]

```
[164]: train.dtypes
```

```
[164]: BMI                                float64
PhysicalHealth                           int64
MentalHealth                             int64
SleepTime                                int64
Smoking_Yes                              uint8
AlcoholDrinking_Yes                      uint8
Stroke_Yes                               uint8
DiffWalking_Yes                          uint8
Sex_Female                               uint8
Sex_Male                                  uint8
AgeCategory_18_24                         uint8
AgeCategory_25_29                         uint8
AgeCategory_30_34                         uint8
AgeCategory_35_39                         uint8
AgeCategory_40_44                         uint8
AgeCategory_45_49                         uint8
AgeCategory_50_54                         uint8
AgeCategory_55_59                         uint8
AgeCategory_60_64                         uint8
AgeCategory_65_69                         uint8
AgeCategory_70_74                         uint8
AgeCategory_75_79                         uint8
AgeCategory_80_or_older                   uint8
```



```

Race_American_Indian_Alaskan_Native    uint8
Race_Asian                              uint8
Race_Black                              uint8
Race_Hispanic                           uint8
Race_Other                              uint8
Race_White                              uint8
Diabetic_No_borderline_diabetes          uint8
Diabetic_Yes                             uint8
Diabetic_Yes_during_pregnancy            uint8
PhysicalActivity_Yes                     uint8
GenHealth_Excellent                      uint8
GenHealth_Fair                           uint8
GenHealth_Good                           uint8
GenHealth_Poor                           uint8
GenHealth_Very_good                      uint8
Asthma_Yes                              uint8
KidneyDisease_Yes                        uint8
SkinCancer_Yes                           uint8
HeartDisease_Yes                         uint8
dtype: object

```

```
[165]: test.columns
```

```

[165]: Index(['BMI', 'PhysicalHealth', 'MentalHealth', 'SleepTime', 'Smoking_Yes',
            'AlcoholDrinking_Yes', 'Stroke_Yes', 'DiffWalking_Yes', 'Sex_Female',
            'Sex_Male', 'AgeCategory_18_24', 'AgeCategory_25_29',
            'AgeCategory_30_34', 'AgeCategory_35_39', 'AgeCategory_40_44',
            'AgeCategory_45_49', 'AgeCategory_50_54', 'AgeCategory_55_59',
            'AgeCategory_60_64', 'AgeCategory_65_69', 'AgeCategory_70_74',
            'AgeCategory_75_79', 'AgeCategory_80_or_older',
            'Race_American_Indian_Alaskan_Native', 'Race_Asian', 'Race_Black',
            'Race_Hispanic', 'Race_Other', 'Race_White',
            'Diabetic_No_borderline_diabetes', 'Diabetic_Yes',
            'Diabetic_Yes_during_pregnancy', 'PhysicalActivity_Yes',
            'GenHealth_Excellent', 'GenHealth_Fair', 'GenHealth_Good',
            'GenHealth_Poor', 'GenHealth_Very_good', 'Asthma_Yes',
            'KidneyDisease_Yes', 'SkinCancer_Yes', 'HeartDisease_Yes'],
            dtype='object')

```

```

[166]: formula = ols_formula(train, "HeartDisease_Yes")
log_model = smf.logit(formula = formula, data = train).fit()
log_model.summary()

```

```

Warning: Maximum number of iterations has been exceeded.
Current function value: 0.227795
Iterations: 35

```

```
C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\base\model.py:566:
```

ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals

warnings.warn("Maximum Likelihood optimization failed to "

[166]: <class 'statsmodels.iolib.summary.Summary'>

"""

Logit Regression Results

```

=====
Dep. Variable:          HeartDisease_Yes    No. Observations:          255836
Model:                  Logit              Df Residuals:          255798
Method:                 MLE                Df Model:              37
Date:                  Mon, 16 May 2022    Pseudo R-squ.:           0.2227
Time:                  18:16:41           Log-Likelihood:         -58278.
converged:              False              LL-Null:                -74974.
Covariance Type:        nonrobust          LLR p-value:             0.000
=====

```

```

=====
                                coef    std err          z      P>|z|
-----
[0.025    0.975]
-----
Intercept                    -2.5441         nan         nan         nan
nan              nan
BMI                          0.0092         0.001         7.203         0.000
0.007          0.012
PhysicalHealth               0.0031         0.001         3.213         0.001
0.001          0.005
MentalHealth                 0.0047         0.001         4.753         0.000
0.003          0.007
SleepTime                   -0.0266         0.005        -5.484         0.000
-0.036        -0.017
Smoking_Yes                  0.3518         0.016        21.952         0.000
0.320          0.383
AlcoholDrinking_Yes         -0.2325         0.037        -6.232         0.000
-0.306        -0.159
Stroke_Yes                   1.0674         0.025        42.298         0.000
1.018          1.117
DiffWalking_Yes              0.2128         0.020        10.498         0.000
0.173          0.253
Sex_Female                  -0.2140         nan         nan         nan
nan              nan
Sex_Male                     0.4948         nan         nan         nan
nan              nan
AgeCategory_18_24            -1.7655        6.69e+05    -2.64e-06         1.000
-1.31e+06      1.31e+06
AgeCategory_25_29            -1.5392        6.7e+05    -2.3e-06         1.000
-1.31e+06      1.31e+06

```

AgeCategory_30_34	-1.2166	6.7e+05	-1.82e-06	1.000
-1.31e+06 1.31e+06				
AgeCategory_35_39	-1.1732	6.72e+05	-1.75e-06	1.000
-1.32e+06 1.32e+06				
AgeCategory_40_44	-0.7311	6.69e+05	-1.09e-06	1.000
-1.31e+06 1.31e+06				
AgeCategory_45_49	-0.4004	6.71e+05	-5.96e-07	1.000
-1.32e+06 1.32e+06				
AgeCategory_50_54	-0.0008	6.7e+05	-1.26e-09	1.000
-1.31e+06 1.31e+06				
AgeCategory_55_59	0.2437	6.71e+05	3.63e-07	1.000
-1.31e+06 1.31e+06				
AgeCategory_60_64	0.5221	6.69e+05	7.81e-07	1.000
-1.31e+06 1.31e+06				
AgeCategory_65_69	0.7508	6.71e+05	1.12e-06	1.000
-1.32e+06 1.32e+06				
AgeCategory_70_74	1.0451	6.72e+05	1.55e-06	1.000
-1.32e+06 1.32e+06				
AgeCategory_75_79	1.2395	6.7e+05	1.85e-06	1.000
-1.31e+06 1.31e+06				
AgeCategory_80_or_older	1.4998	6.71e+05	2.24e-06	1.000
-1.32e+06 1.32e+06				
Race_American_Indian_Alaskan_Native	-0.4505	nan	nan	nan
nan nan				
Race_Asian	-0.9895	nan	nan	nan
nan nan				
Race_Black	-0.8017	nan	nan	nan
nan nan				
Race_Hispanic	-0.7244	nan	nan	nan
nan nan				
Race_Other	-0.5281	nan	nan	nan
nan nan				
Race_White	-0.5436	nan	nan	nan
nan nan				
Diabetic_No_borderline_diabetes	0.1157	0.047	2.472	0.013
0.024 0.207				
Diabetic_Yes	0.4724	0.019	25.299	0.000
0.436 0.509				
Diabetic_Yes_during_pregnancy	0.0881	0.120	0.737	0.461
-0.146 0.322				
PhysicalActivity_Yes	0.0164	0.018	0.913	0.361
-0.019 0.052				
GenHealth_Excellent	-1.3443	nan	nan	nan
nan nan				
GenHealth_Fair	0.1536	nan	nan	nan
nan nan				
GenHealth_Good	-0.3179	nan	nan	nan

```

nan          nan
GenHealth_Poor          0.5423          nan          nan          nan
nan          nan
GenHealth_Very_good    -0.8657          nan          nan          nan
nan          nan
Asthma_Yes             0.2862          0.021          13.358          0.000
0.244          0.328
KidneyDisease_Yes      0.5485          0.027          20.023          0.000
0.495          0.602
SkinCancer_Yes         0.1183          0.022          5.443          0.000
0.076          0.161
=====
=====
"""


```

Questions - Why are GenHealth, Race, Sex P-values equal to nan? perfect multicollinearity - exclude one of the dummy variables

```

[167]: ypred = log_model.predict(train)
p, r, thresholds = precision_recall_curve(y_train, ypred)
plot_precision_recall_vs_threshold(p, r, thresholds)

```



```
[168]: confusion_matrix_train(log_model, 0.2)
```

```
Accuracy = 87.50762207038885
Precision = 33.84420970773119
FNR = 52.49681644533382
TPR or Recall = 47.50318355466618
Confusion matrix =
      Predicted 0 Predicted 1
Actual 0      213431.0      20417.0
Actual 1       11543.0      10445.0
```

```
[168]: ' '
```

```
[169]: confusion_matrix_test(test,test.HeartDisease_Yes,log_model, 0.2)
```

```
Accuracy = 87.59830516424584
Precision = 33.696069645371914
FNR = 51.12349117920149
TPR or Recall = 48.87650882079851
Confusion matrix =
      Predicted 0 Predicted 1
Actual 0      213431.0      20417.0
Actual 1      11543.0      10445.0
```

```
[169]: ' '
```

```
[ ]:
```

6 Classification Tree

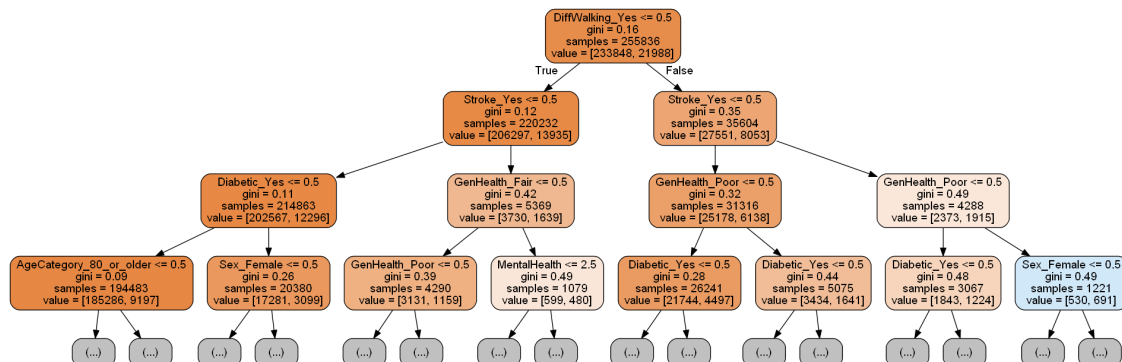
```
[175]: #Defining the object to build a regression tree
model = DecisionTreeClassifier(random_state=1, max_depth=5)

#Fitting the regression tree to the data
model.fit(X_train, y_train)
```

```
[175]: DecisionTreeClassifier(max_depth=5, random_state=1)
```

```
[176]: #Visualizing the regression tree
dot_data = StringIO()
export_graphviz(model, out_file=dot_data,
                filled=True, rounded=True, max_depth = 3,
                feature_names = X_train.columns,precision=2)
graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
#graph.write_png('car_price_tree.png')
Image(graph.create_png())
```

```
[176]:
```



```
[177]: confusion_matrix_data(X_train,y_train,model,cutoff=0.4)
```

```
Accuracy = 91.41129473569005
Precision = 50.149730485126774
FNR = 88.57558668364562
TPR or Recall = 11.424413316354375
Confusion matrix =
      Predicted 0 Predicted 1
Actual 0      231351.0      2497.0
Actual 1      19476.0      2512.0
```

```
[177]: ' '
```

```
[178]: X_train.columns
```

```
[178]: Index(['BMI', 'PhysicalHealth', 'MentalHealth', 'SleepTime', 'Smoking_Yes',
            'AlcoholDrinking_Yes', 'Stroke_Yes', 'DiffWalking_Yes', 'Sex_Female',
            'Sex_Male', 'AgeCategory_18_24', 'AgeCategory_25_29',
            'AgeCategory_30_34', 'AgeCategory_35_39', 'AgeCategory_40_44',
            'AgeCategory_45_49', 'AgeCategory_50_54', 'AgeCategory_55_59',
            'AgeCategory_60_64', 'AgeCategory_65_69', 'AgeCategory_70_74',
            'AgeCategory_75_79', 'AgeCategory_80_or_older',
            'Race_American_Indian_Alaskan_Native', 'Race_Asian', 'Race_Black',
            'Race_Hispanic', 'Race_Other', 'Race_White',
            'Diabetic_No_borderline_diabetes', 'Diabetic_Yes',
            'Diabetic_Yes_during_pregnancy', 'PhysicalActivity_Yes',
            'GenHealth_Excellent', 'GenHealth_Fair', 'GenHealth_Good',
            'GenHealth_Poor', 'GenHealth_Very_good', 'Asthma_Yes',
            'KidneyDisease_Yes', 'SkinCancer_Yes'],
            dtype='object')
```

```
[179]: #Defining parameters and the range of values over which to optimize
param_grid = {
    'max_depth': range(1,10),
    'max_leaf_nodes': range(10,30),
}

skf = StratifiedKFold(n_splits=5)
#The folds are made by preserving the percentage of samples for each class.
#Minimizing FNR is equivalent to maximizing recall
grid_search = GridSearchCV(DecisionTreeClassifier(random_state=1), param_grid,
                           scoring=['precision','recall','accuracy'],
                           refit="recall", cv=skf, n_jobs=-1, verbose = True)
grid_search.fit(X_train.iloc[:10000, :], y_train.iloc[:10000])
# make the predictions
y_pred = grid_search.predict(X_test)
print('Best params for recall')
```

```
print(grid_search.best_params_)
```

Fitting 5 folds for each of 180 candidates, totalling 900 fits

Best params for recall

```
{'max_depth': 5, 'max_leaf_nodes': 27}
```

```
[180]: #All results of the grid search can be seen with cv_results_  
cv_scores = pd.DataFrame(grid_search.cv_results_)  
  
model = DecisionTreeClassifier(random_state=1, max_depth = 5, max_leaf_nodes=27)  
model.fit(X_train,y_train)  
print(confusion_matrix_data(X_train,y_train,model))  
print(confusion_matrix_data(X_test,y_test,model))
```

Accuracy = 91.50823183601995

Precision = 59.66201322556943

FNR = 96.3070765872294

TPR or Recall = 3.692923412770602

Confusion matrix =

	Predicted 0	Predicted 1
Actual 0	233299.0	549.0
Actual 1	21176.0	812.0

Accuracy = 91.60868681499085

Precision = 52.71084337349398


FNR = 96.7502321262767

TPR or Recall = 3.2497678737233056

Confusion matrix =

	Predicted 0	Predicted 1
Actual 0	58417.0	157.0
Actual 1	5210.0	175.0

```
[181]: ypred = model.predict_proba(X_train)[:, 1]  
p, r, thresholds = precision_recall_curve(y_train, ypred)  
plot_precision_recall_vs_threshold(p, r, thresholds)
```

```
[182]: print(confusion_matrix_data(X_train,y_train,model,cutoff=0.2))
       print(confusion_matrix_data(X_test,y_test,model,cutoff=0.2))
```

```
Accuracy = 85.8202129489204
Precision = 29.005906379477537
FNR = 55.10733127160269
TPR or Recall = 44.89266872839731
Confusion matrix =
      Predicted 0 Predicted 1
Actual 0      209688.0      24160.0
Actual 1      12117.0       9871.0

Accuracy = 86.04574805734924
```

```
Precision = 28.90345649582837
FNR = 54.967502321262764
TPR or Recall = 45.032497678737236
Confusion matrix =
      Predicted 0 Predicted 1
Actual 0      52609.0      5965.0
Actual 1      2960.0      2425.0
```

Cost Complexity Pruning

```
[183]: model = DecisionTreeClassifier(random_state = 1) #model without any restrictions
path= model.cost_complexity_pruning_path(X_train.iloc[:10000, :],y_train.iloc[:
→10000]) # Compute the pruning path during Minimal Cost-Complexity Pruning.
```

```
[184]: alphas=path['ccp_alphas']
len(alphas)
```

```
[184]: 382
```

```
[185]: #Grid search to optimize parameter values

skf = StratifiedKFold(n_splits=5)
grid_search = GridSearchCV(DecisionTreeClassifier(random_state = 1), param_grid_
→= {'ccp_alpha':alphas},

→scoring=['precision','recall','accuracy'],
refit="recall", cv=skf,
→n_jobs=-1, verbose = True)
grid_search.fit(X_train.iloc[:10000, :], y_train.iloc[:10000])

# make the predictions
y_pred = grid_search.predict(X_test)

print('Best params for recall')
print(grid_search.best_params_)
```

```
Fitting 5 folds for each of 382 candidates, totalling 1910 fits
Best params for recall
{'ccp_alpha': 0.0}
```


```
[186]: tree = DecisionTreeClassifier(ccp_alpha=0.0,random_state=1)
tree.fit(X_train, y_train)
print(confusion_matrix_data(X_train,y_train,tree,cutoff = 0.2))
print(confusion_matrix_data(X_test,y_test,tree, cutoff = 0.2))
```

```
Accuracy = 99.52899513750997
Precision = 94.88187275946962
```

```
FNR = 0.09095870474804439
TPR or Recall = 99.90904129525195
Confusion matrix =
      Predicted 0 Predicted 1
Actual 0      232663.0      1185.0
Actual 1         20.0      21968.0
```

```
Accuracy = 86.06607357838615
Precision = 22.074425969912905
FNR = 74.11327762302693
TPR or Recall = 25.886722376973072
Confusion matrix =
      Predicted 0 Predicted 1
Actual 0      53653.0      4921.0
Actual 1      3991.0      1394.0
```

```
[187]: ypred = tree.predict_proba(X_test)[:, 1]
p, r, thresholds = precision_recall_curve(y_test, ypred)
plot_precision_recall_vs_threshold(p, r, thresholds)
```



[]:

7 Random forests

Categorical Age

```
[314]: recall = [0]*5
```

```
[315]: i = 0
start_time = time.time()

for pr in range(10, 15):
```

```

model = RandomForestClassifier(random_state = 1, oob_score = True, verbose_
↪ = False,
                                n_estimators = 500, max_features = pr,
                                n_jobs = -1).fit(X_train, y_train)
oob_pred = model.oob_decision_function_[:, 1]
bins = np.array([0, 0.5, 1])
cm = np.histogram2d(y_train, oob_pred, bins = bins)[0]
recall[i] = 100*(cm[1, 1]) / (cm[1, 0] + cm[1, 1])
i += 1

end_time = time.time()

#took 0.2 min for range of 5 and 4000 obs
print('Time taken = ', (end_time-start_time)/60, 'min')
print('max recall = ', np.max(recall) + 10)
print('Best value of max_features = ', np.argmax(recall) + 10)

```

Time taken = 0.18982399702072145 min
max recall = 87.65
Best value of max_features = 12

```

[318]: model_rf = RandomForestClassifier(random_state = 1, n_jobs = -1, max_features = _
↪ 12,
                                n_estimators = 500).fit(X_train, y_train)

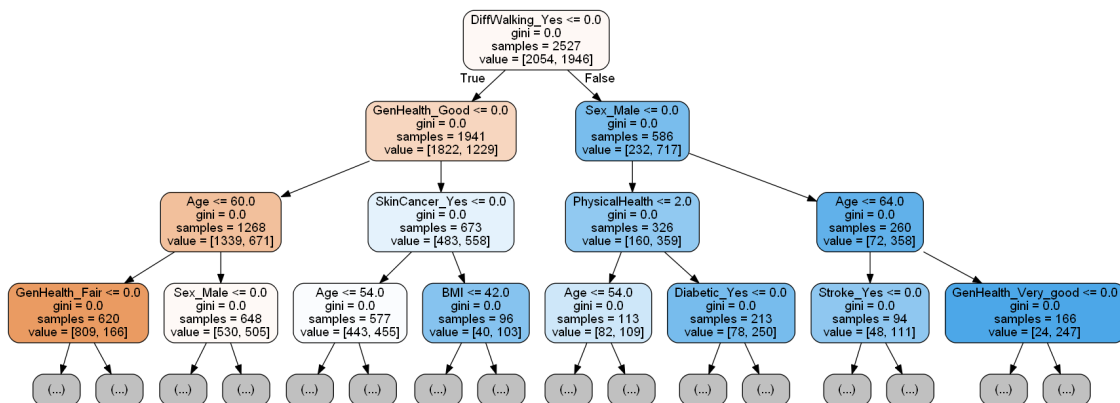
```

```


[320]: dot_data = StringIO()
export_graphviz(model_rf.estimators_[0], out_file = dot_data,
                filled = True, rounded = True, max_depth = 3,
                feature_names = X_train.columns, precision=0)
graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
#graph.write_png('car_price_tree.png')
Image(graph.create_png())

```

[320]:



```
[321]: ypred = model_rf.predict_proba(X_test)[: , 1]
p, r, thresholds = precision_recall_curve(y_test, ypred)
plot_precision_recall_vs_threshold(p, r, thresholds)
```



```
[327]: confusion_matrix_data(X_train, y_train, model_rf, 0.4)
```

```
Accuracy = 99.875
Precision = 99.75062344139651
FNR = 0.0
TPR or Recall = 100.0
Confusion matrix =
      Predicted 0 Predicted 1
Actual 0      1995.0       5.0
```

```
Actual 1          0.0        2000.0
```

```
[327]: ' '
```

```
[334]: confusion_matrix_data(X_test, y_test, model_rf, 0.4)
```

```
Accuracy = 70.2
Precision = 65.93059936908517
FNR = 16.4
TPR or Recall = 83.6
Confusion matrix =
      Predicted 0 Predicted 1
Actual 0      284.0      216.0
Actual 1      82.0      418.0
```


```
[334]: ' '
```

- Using cutoff = 0.4, accuracy decreases 0.5%, precision decreases 3%, but recall improves by 8%

Using X_train2

```
[335]: model_rf2 = RandomForestClassifier(random_state = 1, n_jobs = -1, max_features_
      ↪= 12,
      n_estimators = 500).fit(X_train2, y_train2)
```

```
[341]: ypred = model_rf2.predict_proba(X_test)[: , 1]
p, r, thresholds = precision_recall_curve(y_test, ypred)
plot_precision_recall_vs_threshold(p, r, thresholds)
```



```
[350]: confusion_matrix_data(X_train2, y_train2, model_rf2, 0.5)
```

```
Accuracy = 99.775
Precision = 99.79989994997499
FNR = 0.25
TPR or Recall = 99.75
Confusion matrix =
      Predicted 0 Predicted 1
Actual 0      1996.0         4.0
Actual 1         5.0      1995.0
```

```
[350]: ' '
```



```
[349]: confusion_matrix_data(X_test, y_test, model_rf2, 0.5)
```

```
Accuracy = 76.3
Precision = 74.57943925233644
FNR = 20.2
TPR or Recall = 79.8
Confusion matrix =
      Predicted 0 Predicted 1
Actual 0      364.0      136.0
Actual 1      101.0      399.0
```

```
[349]: ' '
```

8 AdaBoost

```
[44]: #tuning params
start_time = time.time()

model_ab = AdaBoostClassifier(random_state = 1)
grid = dict()
grid['n_estimators'] = [10, 50, 100, 200, 500]
grid['learning_rate'] = [0.0001, 0.001, 0.01, 0.1, 1.0]
grid['base_estimator'] = [DecisionTreeClassifier(max_depth = 1),
    ↳ DecisionTreeClassifier(max_depth = 2),
    ↳ DecisionTreeClassifier(max_depth = 3),
    ↳ DecisionTreeClassifier(max_depth = 4)]

cv = StratifiedKFold(n_splits = 5, shuffle = True, random_state = 1)
grid_search = GridSearchCV(estimator = model_ab, param_grid = grid, n_jobs = -1,
    cv = cv, scoring = ['precision', 'recall',
    ↳ 'accuracy'], refit = 'recall')
grid_result = grid_search.fit(X_train.iloc[:1000], y_train.iloc[:1000])


print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
print("Time taken = ", (time.time()-start_time)/60, " minutes")
```

```
Best: 0.223529 using {'base_estimator': DecisionTreeClassifier(max_depth=2),
'learning_rate': 1.0, 'n_estimators': 100}
Time taken = 1.7052798986434936 minutes
```

- took 1.7 minutes with 1000 obs

```
[46]: #creating model with optimal params
model_ab = AdaBoostClassifier(random_state = 1, base_estimator =
    ↳ DecisionTreeClassifier(max_depth = 2),
    ↳ learning_rate = 1.0, n_estimators = 100).
    ↳ fit(X_train.iloc[:10000], y_train.iloc[:10000])
```

```
[47]: ypred = model_ab.predict_proba(X_test)[: , 1]
      p, r, thresholds = precision_recall_curve(y_test, ypred)
      plot_precision_recall_vs_threshold(p, r, thresholds)
```



```
[49]: confusion_matrix_data(X_train, y_train, model_ab, 0.5)
```

```
Accuracy = 90.75579668224957
Precision = 38.88740304894357
FNR = 86.77460432963434
TPR or Recall = 13.225395670365653
Confusion matrix =
      Predicted 0 Predicted 1
Actual 0      229278.0      4570.0
```

Actual 1	19080.0	2908.0
----------	---------	--------

[49]: ' '

```
[51]: confusion_matrix_data(X_test, y_test, model_ab, 0.5)
```

```
Accuracy = 90.88947607060773
Precision = 38.29449152542373
FNR = 86.57381615598885
TPR or Recall = 13.426183844011142
Confusion matrix =
      Predicted 0 Predicted 1
Actual 0      57409.0      1165.0
Actual 1      4662.0       723.0
```

[51]: ' '

[]:

9 Gradient Boosting

```
[52]: #tuning params
start_time = time.time()

model_gb = GradientBoostingClassifier(random_state = 1, max_features = 'sqrt')
grid = dict()
grid['n_estimators'] = [10, 50, 100, 200, 500]
grid['learning_rate'] = [0.0001, 0.001, 0.01, 0.1, 1.0]
grid['max_depth'] = [1, 2, 3, 4, 5]
grid['subsample'] = [0.5, 1.0]

cv = StratifiedKFold(n_splits = 5, shuffle = True, random_state = 1)
grid_search = GridSearchCV(estimator = model_gb, param_grid = grid, n_jobs = -1,
                           cv = cv, scoring = ['precision', 'recall',
→ 'accuracy'],
                           refit = 'recall')
grid_result = grid_search.fit(X_train.iloc[:1000], y_train.iloc[:1000])

print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
print("Time taken = ", (time.time()-start_time)/60, " minutes")
```


```
Best: 0.458824 using {'learning_rate': 1.0, 'max_depth': 1, 'n_estimators': 100,
'subsample': 0.5}
```

```
Time taken = 1.3608951171239216 minutes
```

- took 1.36 min using 1000 obs

```
[53]: model_gb = GradientBoostingClassifier(random_state = 1, max_depth = 1,
→learning_rate = 1.0,
n_estimators = 100, subsample = 0.5).
→fit(X_train.iloc[:10000], y_train.iloc[:10000])
```

```
[54]: ypred = model_gb.predict_proba(X_test)[: , 1]
p, r, thresholds = precision_recall_curve(y_test, ypred)
plot_precision_recall_vs_threshold(p, r, thresholds)
```



```
[58]: confusion_matrix_data(X_train, y_train, model_gb, 0.2)
```

Accuracy = 87.33211901374318
Precision = 32.378504514863536

```

FNR = 56.458068037111154
TPR or Recall = 43.541931962888846
Confusion matrix =
      Predicted 0 Predicted 1
Actual 0      213853.0      19995.0
Actual 1      12414.0       9574.0

```

[58]: ' '

[59]: `confusion_matrix_data(X_test, y_test, model_gb, 0.2)`

```

Accuracy = 87.62957519661033
Precision = 32.92336802270577
FNR = 54.76323119777159
TPR or Recall = 45.23676880222841
Confusion matrix =
      Predicted 0 Predicted 1
Actual 0      53611.0      4963.0
Actual 1       2949.0      2436.0

```

[59]: ' '